

# Pointers vs Values: digging into the performance war

Mario Macías Lloret

Senior software engineer at New Relic



Barcelona Golang Meetup  
January 2020



[twitter.com/MaciasUPC](https://twitter.com/MaciasUPC)

[github.com/mariomac](https://github.com/mariomac)

# Who am I?




# The Core Agents and Open Standards Team

JF



Mario



Juan



Antonio




WE NEED YOU!



Carlos



David



Cristian



Frank

# Table of contents

## **Refreshing: values vs pointers**

Digging more:  $\mu$ Benchmarks

A more realistic use case

Let New Relic measure it!

Conclusions

# Using values vs using pointers

- Values

- Safe against nil
- Cleaner
  - no need to check for nil
  - no pointer operators \*, &

- Pointers

- Allow passing arguments by reference
- Allow sharing a common state between different instances

The super-optimizer opinion



*“We should forget about small efficiencies, say about 97% of the time: premature optimization is the root of all evil”*


-- Donald Knuth



# January 2019 BCN Golang Meetup

- T(Heap allocation) >> T(Stack Allocation)
- Golang applies “Escape Analysis” techniques to infer where an object is allocated
- Abuse of pointers escape values to the heap

```
func a() *Obj {  
    r := Obj{}  
    // ... do something  
    return &r;  
}  
  
func b() {  
    o := a()  
    // ... do something  
}
```







# January 2019 BCN Golang Meetup

- T(Heap allocation) >> T(Stack Allocation)

- **Benchmark Results**

- `go test ./donut/. -bench=Benchmark -benchmem`

f	BenchmarkValue-4	5000000	262 ns/op	15 B/op	0 allocs/op
	BenchmarkPointers-4	5000000	332 ns/op	79 B/op	1 allocs/op

Our code (including random number generation and scoring operations) using values is  
~23% faster than using pointers!

```
}  
func b() {  
    o := a()  
    // ... do something  
}
```

# Table of contents

Refreshing: values vs pointers

## **Digging more: $\mu$ Benchmarks**

A more realistic use case

Let New Relic measure it!

Conclusions

# Digging more: $\mu$ Benchmarks



- Small, localized benchmarks to test a single system functionality.
- Not really meaningful from a wider application point of view

```
type Foo struct {  
    A int  
    B int  
    C int  
}
```

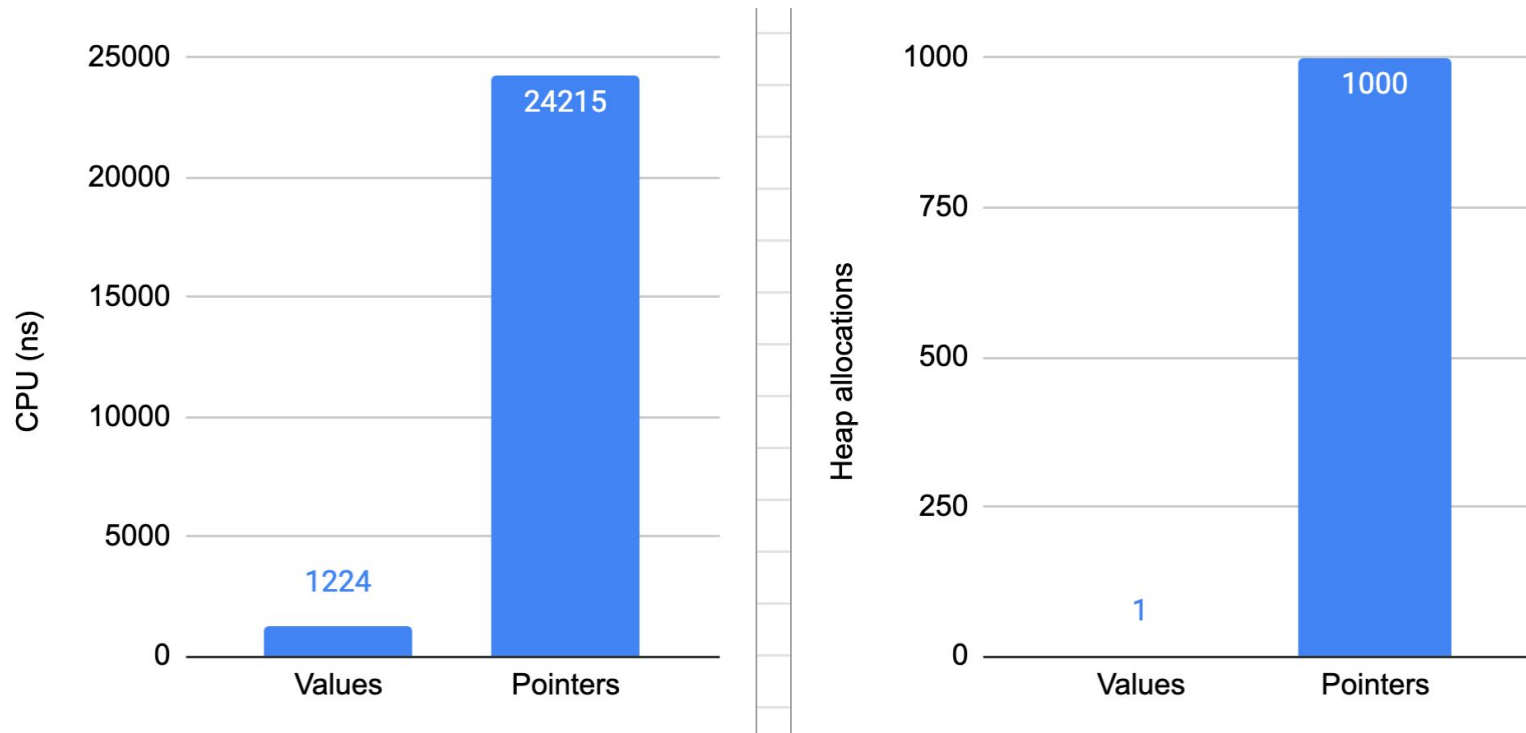
```
const FooLength = 1000  
  
func addFoos(foos []Foo) []Foo {  
    func addFoosP(foos []*Foo) []*Foo {  
        for i := 0; i < FooLength; i++ {  
            foos = append(foos, &Foo{  
                A: i,  
            })  
        }  
    }  
    return foos  
}
```

# μBenchmarks: initial results

```
func BenchmarkSliceCreation_Values(b *testing.B) {  
    for i := 0; i < b.N; i++ {  
        // ...  
    }  
}  
  
func BenchmarkSliceCreation_Pointers(b *testing.B) {  
    for i := 0; i < b.N; i++ {  
        // ...  
    }  
}
```

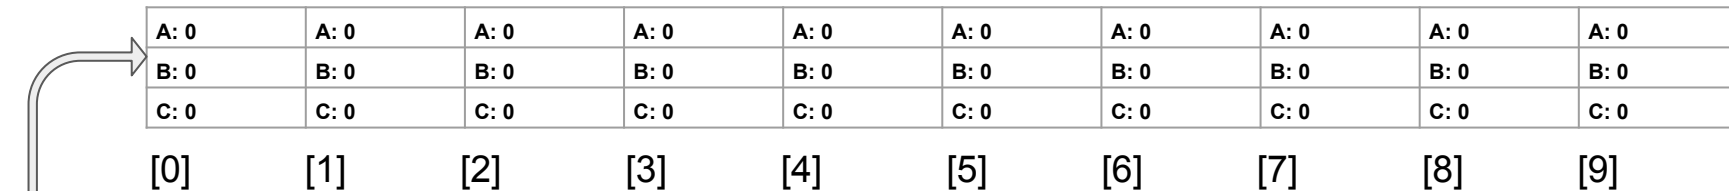


# μBenchmarks: pre-allocated arrays



# Adding a value to an array

```
arr := make([]Foo, 0, 10)
```



A: 0	A: 0	A: 0	A: 0	A: 0	A: 0	A: 0	A: 0	A: 0	A: 0
B: 0	B: 0	B: 0	B: 0	B: 0	B: 0	B: 0	B: 0	B: 0	B: 0
C: 0	C: 0	C: 0	C: 0	C: 0	C: 0	C: 0	C: 0	C: 0	C: 0
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]	[9]

arr

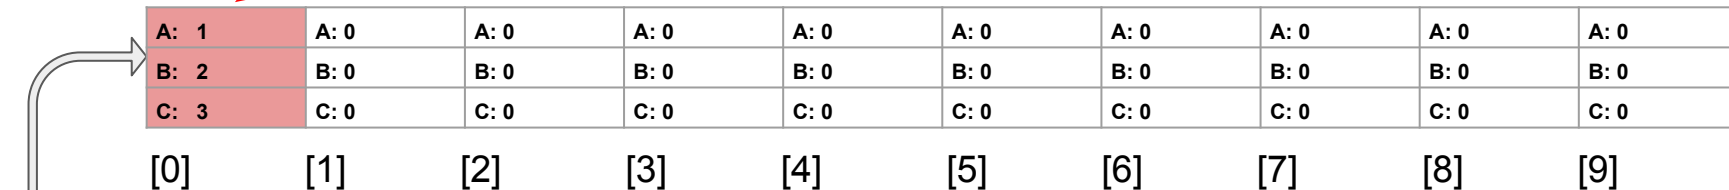
- length: 0
- capacity: 10

# Adding a value to an array

```
arr := make([]Foo, 0, 10)
```

```
arr = append(arr, Foo{A: 1, B: 2, C: 3})
```

copy




A: 1	A: 0	A: 0	A: 0	A: 0	A: 0	A: 0	A: 0	A: 0	A: 0
B: 2	B: 0	B: 0	B: 0	B: 0	B: 0	B: 0	B: 0	B: 0	B: 0
C: 3	C: 0	C: 0	C: 0	C: 0	C: 0	C: 0	C: 0	C: 0	C: 0
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]	[9]

arr

- length: **1**
- capacity: 10

# Adding a reference to an array

```
arr := make([]*Foo, 0, 10)
```



arr

- length: 0
- capacity: 10




# Adding a reference to an array

```
arr := make([]*Foo, 0, 10)  
arr = append(arr, &Foo{A: 1, B: 2, C: 3})
```

A: 1
B: 2
C: 3

Heap allocation

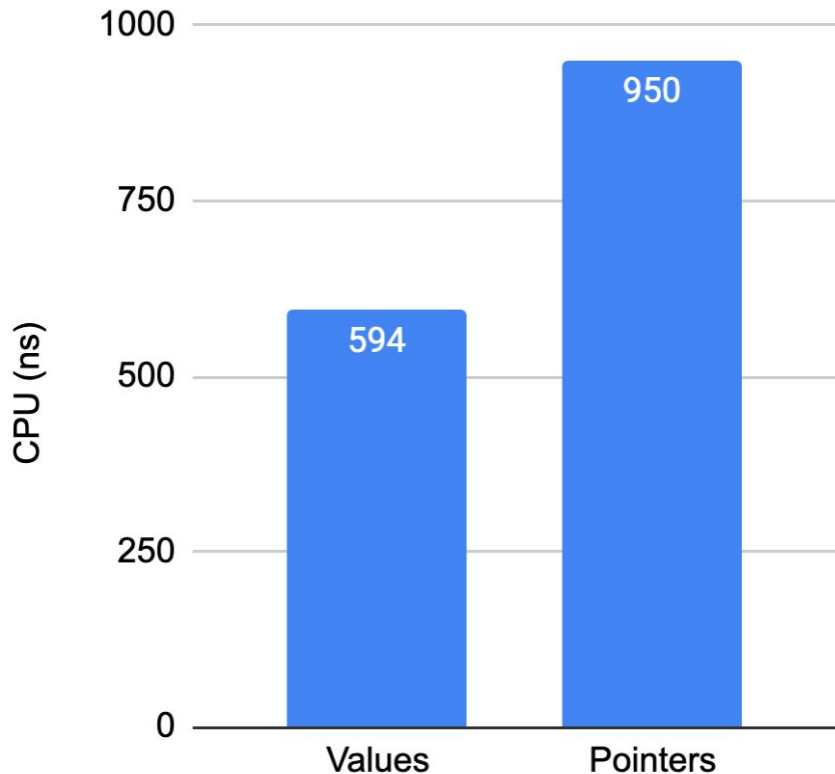


arr

- length: 1
- capacity: 10

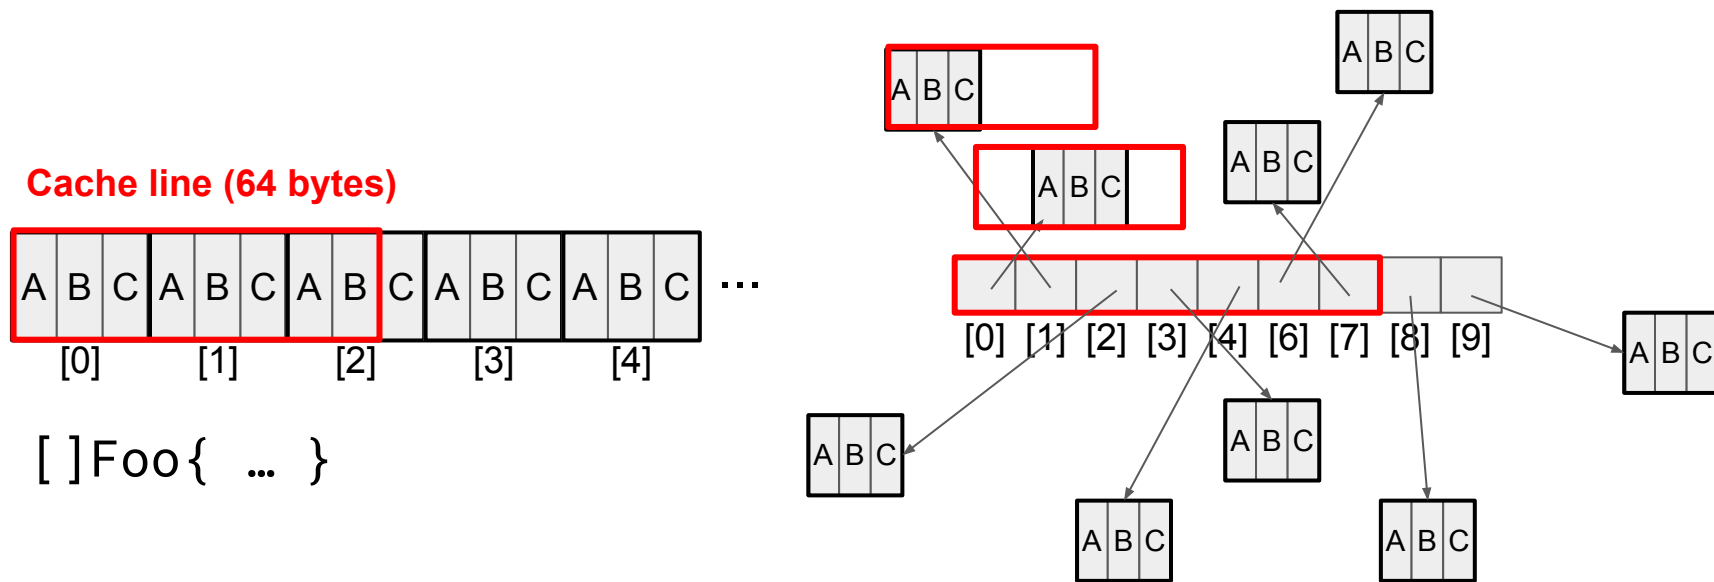
# μBenchmark: array iteration (no allocations)

```
func sumAll(foos []Foo) int {  
    sum := 0  
    for _, f := range foos {  
        sum += f.A + f.B + f.C  
    }  
    return sum  
}
```



```
func sumAll(foos []*Foo) int {  
    for _, f := range foos {  
        f.A + f.B + f.C  
    }  
}
```

# Enforcing cache memory contiguity

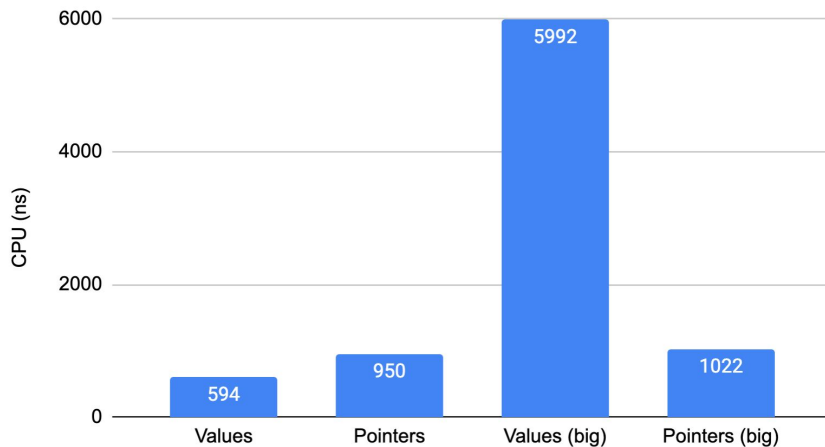


**Increased cache misses**

# μBenchmark: structs >64 bytes

```
type Foo struct {  
    A int  
    B int  
    C int  
    D int  
    E int  
    F int  
    G int  
    H int  
    I int  
    J int  
    K int  
}
```


Slice iteration



# Local μoptimization: minimize local var copy

```
func sumAllLR(foos []Foo) int {  
    sum := 0  
    for i := range foos {  
        f := &foos[i]  
        sum += f.A + f.K  
    }  
    return sum  
}
```

Slice iteration



# Table of contents

Refreshing: values vs pointers

Digging more:  $\mu$ Benchmarks

**A more realistic use case**


Let New Relic measure it!

Conclusions

# Real world bench: dimensional metrics translator

## New relic Flat sample

- event\_type: SystemSample
- operatingSystem: Linux
- agentVersion: 1.8.82
- cpuPercent: 30
- diskFreePercent: 85
- hostname: ip-AC1F0D60
- instanceType: t2.small
- memoryUsedBytes: 1109519701
- etc....



## Dimensional metrics sample

### Common:

- event\_type: SystemSample
- operatingSystem: Linux
- agentVersion: 1.8.82
- hostname: ip-AC1F0D60
- instanceType: t2.small

### Metrics:

name: cpuPercent  
value: 30

name: diskFreePercent  
value: 85

name: memoryUsedBytes  
value: 1109519701

# Dimensional metrics translator

```
{  
  "event_type": "SystemSample",  
  "operatingSystem": "Linux",  
  "agentVersion": "1.8.82",  
  "cpuPercent": 30,  
  "diskFreePercent": 85,  
  "hostname": "ip-AC1F0D60",  
  "instanceType": "t2.small",  
  "memoryUsedBytes": 1109519701  
}
```

json.Unmarshal(..., &map)

dim.FromFlat(map)

submitter.Submit(...)

json.Marshal(...)

```
{  
  "Common": {  
    "event_type": "SystemSample",  
    "operatingSystem": "Linux",  
    "agentVersion": "1.8.82",  
    "hostname": "ip-AC1F0D60",  
    "instanceType": "t2.small"  
  },  
  "Metrics": [{  
    "name": "cpuPercent",  
    "type": "Gauge", "value": 30  
  }, {  
    "name": "diskFreePercent",  
    "type": "Gauge", "value": 85  
  }, {  
    "name": "memoryUsedBytes",  
    "type": "Gauge",  
    "value": 1109519701  
  }]  
}
```



# Benchmark: same code, 2 versions

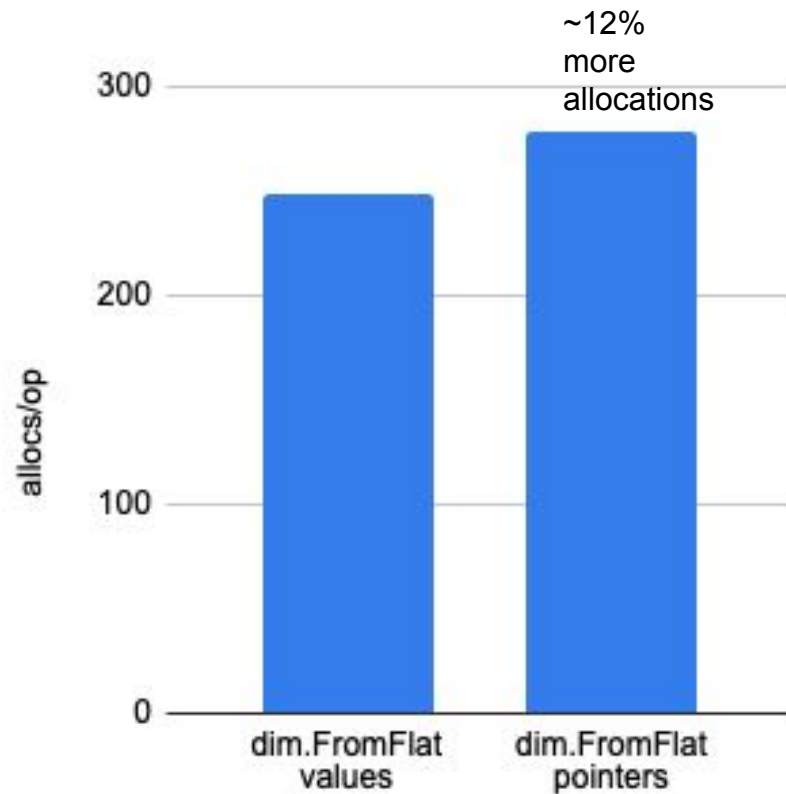
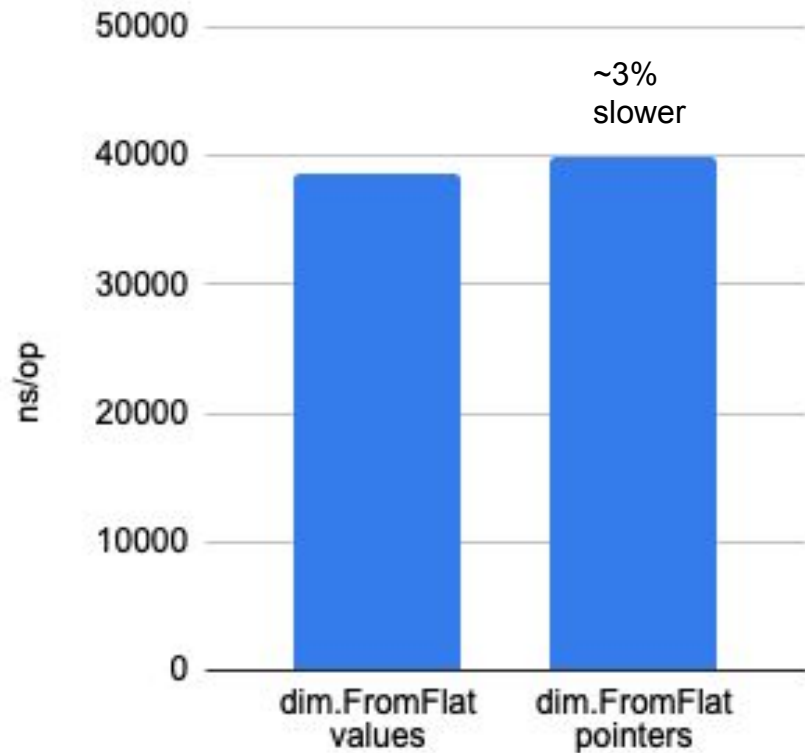
```
type Type string
type Payload struct {
    Common *Common
    Metrics []*Metric
}
type Common struct {
    Attributes map[string]string
    Timestamp int64
}
type Metric struct {
    Name string
    Type Type
    Value float64
}
type Submitter interface {
    Submit(p *Payload) error
}
```

```
func FromFlat(values map[string]interface{}) *Payload
```

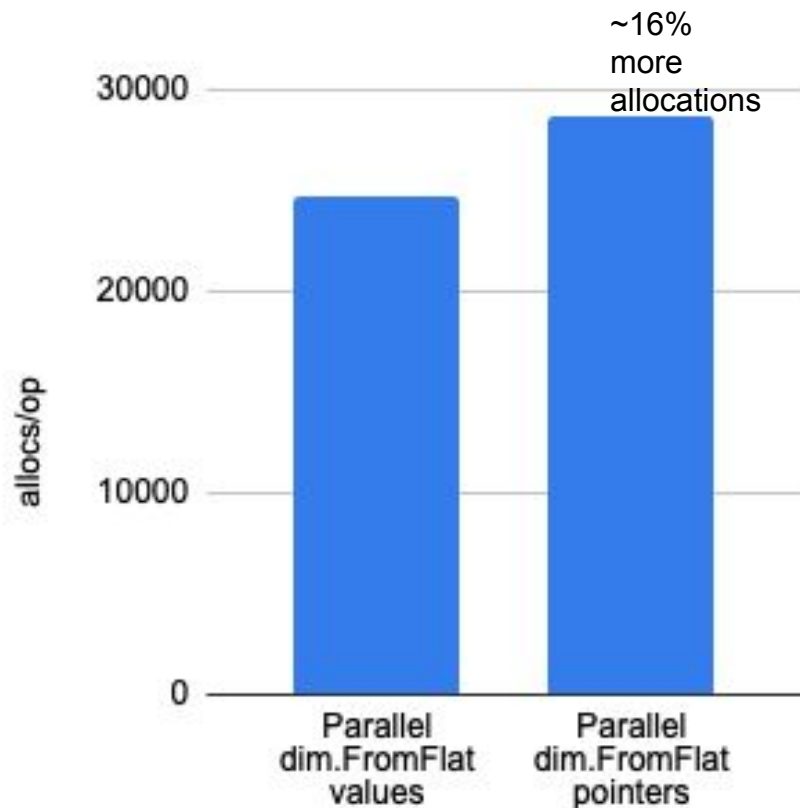
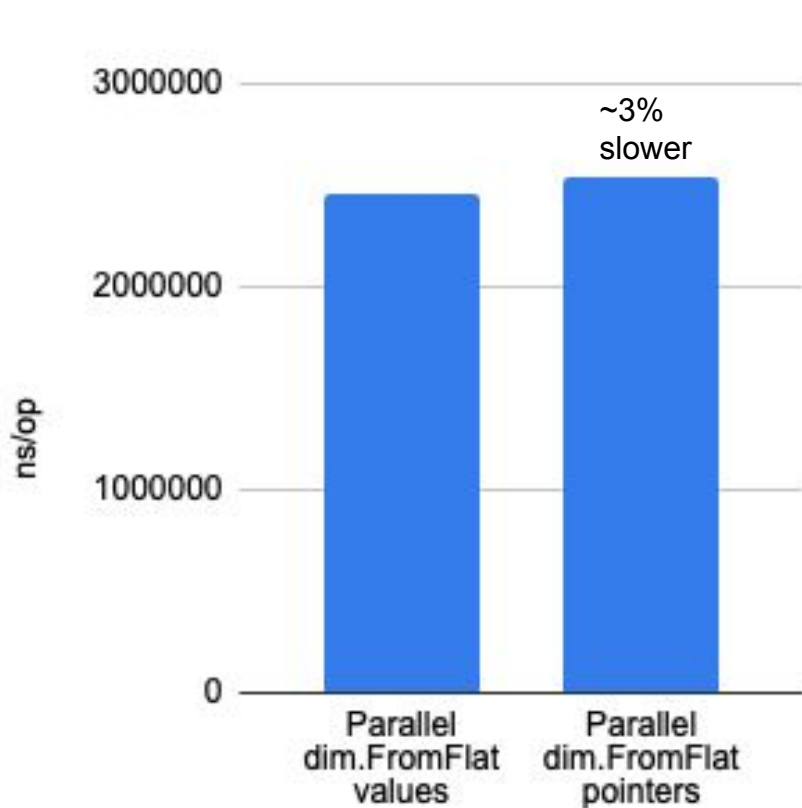
```
type Type string
type Payload struct {
    Common Common
    Metrics []Metric
}
type Common struct {
    Attributes map[string]string
    Timestamp int64
}
type Metric struct {
    Name string
    Type Type
    Value float64
}
type Submitter interface {
    Submit(p Payload) error
}
```

```
func FromFlat(values map[string]interface{}) Payload
```

# Benchmark Results



# Benchmark Results (100 parallel goroutines)



# Table of contents

Refreshing: values vs pointers

Digging more:  $\mu$ Benchmarks

A more realistic use case

**Let New Relic measure it!**

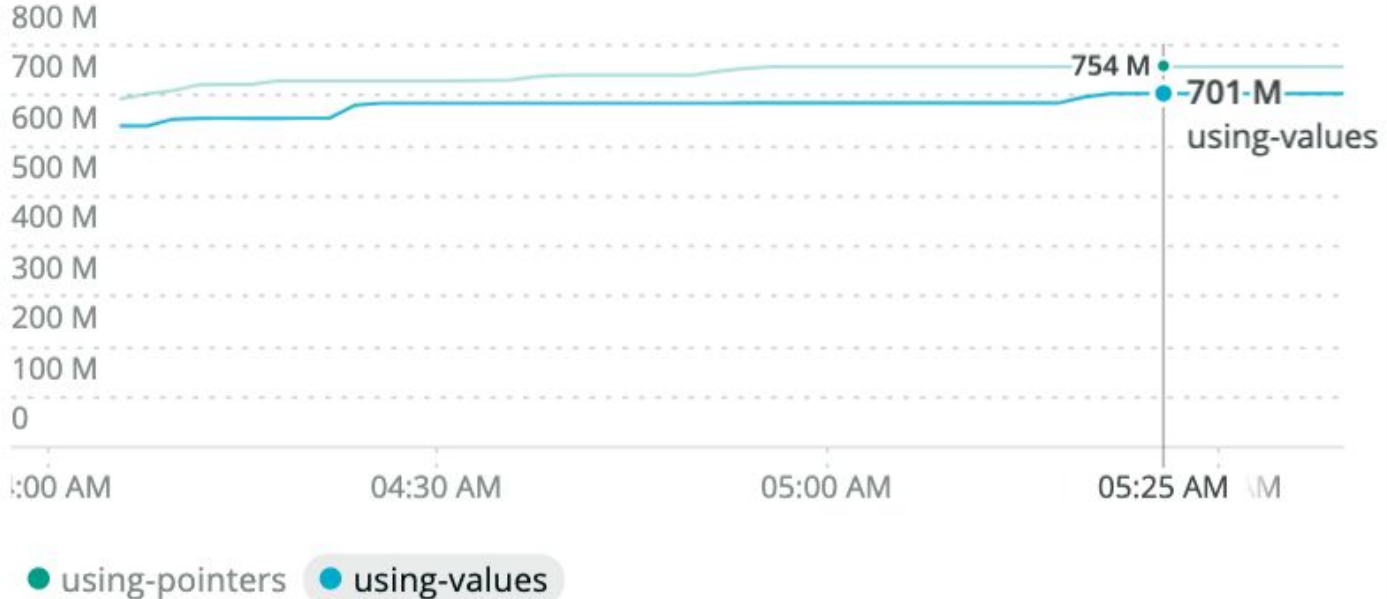
Conclusions

# Running the example in Kubernetes. New Relic Metrics

## K8s Container Memory usage


Since 96 minutes ago

...



# Running the example in Kubernetes. New Relic Metrics

Used by Kubernetes OOM killer



With thousands of containers, this +35% might do the difference

# Table of contents

Refreshing: values vs pointers

Digging more:  $\mu$ Benchmarks

A more realistic use case

Let New Relic measure it!

**Conclusions**

# Conclusions

- First, aim for clean and robust code
- If, in a hot spot, performance is so critical that you must start micro-optimizing:
  - Consider reducing your memory generation (Heap allocations) rather than the memory copy
  - Consider the memory contiguity
  - Consider adding a comment:

```
for i := range foos {  
    f := &foos[i] // DON'T CHANGE THIS!!  
    sum += f.A + f.K  
}
```





# Thank you for your attention!

Mario Macías Lloret  
Senior software engineer at New Relic

Barcelona Golang Meetup  
January 2020

[twitter.com/MaciasUPC](https://twitter.com/MaciasUPC)

[github.com/mariomac](https://github.com/mariomac)